

Implementation and evaluation of Learning Algorithms in Python

HWI - ML Sapienza

Giulio Serra 1904089

November 10, 2019

Contents

1	Introduction	1
2	Understanding the Classification Problem	1
3	Learning Algorithms	1
3.1	Decision Tree	1
3.2	Multinomial Naive Bayes	2
4	Implementations in Python	2
4.1	Parsing the Train Set	2
4.2	Vectorizing the Data	3
4.3	Splitting the DataSet	5
4.4	Creating the Model and Processing the Data	5
5	Evaluating the Performances	7
5.1	Precision	7
5.2	Recall	7
5.3	F1 Score	7
5.4	The Precision-Recall curve	8
6	BlindTest	9
6.1	Binary Classification Problem	9
6.2	Multiclass Classification Problem	11
7	References	12

1 Introduction

This essay will describe the implementation of two learning algorithm using Python and the Sklearn library, evaluating the performances based on the accuracy of the prediction , varying some parameters such as features and dataset.

2 Understanding the Classification Problem

The aim of this implementation is to train two learning algorithms to recognize instructions parsed from a file, by categorizing them by High and Low type and by Compiler.

So it's a problem of text classification: each new instruction need to be categorized and then the model will autonomusly put new words in one of the categories once is trained.

In the fist case, we have a binary classifier(each instruction could be categorize only in high or low), while in the second scenario we have three possible types of instructions: GCC, ICC and CLANG, so we are talking about a multiclass classsifier.

3 Learning Algorithms

A learning algorithm can predict the value of an input set of data within an acceptable range, generally speaking, more data is feed to the algorithm, higher the accuracy of the algorithm increase.

Having two set of data X(input data) and Y(expected values) a learning algorithm can compute a function $h(x)$ (hypothesis function) that is more or less similar to the original function $f(x) = y$ that $h(x)$ tries to mimic.

In order to train the learning algorithm, we have to produce a set of data (rappresenting a sample of all the possible values of the X set), than divide it into two sub-sets:a train set and test set, the train set is used to train the algorithm, to predict the correct values given x inputs.The test set on the other hand, has the purpose to validate the predictions of a learning algorithm once is trained.

The same needs to be done to the Y set that needs to be splitted between test and train set.

3.1 Decision Tree

The Decision Tree is a supervised learning algorithm that creates a tree based on conjunctions of disjunctions that leds to positive results.

Each node of the Decision Tree represents an attribute of the tuple of the dataset, each branch represents a decision (possible value of the tuple), after a certain number of decisions we eventually get an outcome.

3.2 Multinomial Naive Bayes

The multinomial Naive Bayes Classifier is a supervised learning algorithm that belongs to a family of probabilistic classifiers based on the Bayes Rule, it works by assigning labels to the problem instances using vectors of features values, using the maximum likelihood (MLE) by estimating the parameters of a probabilistic distributions.

4 Implementations in Python

The two learning algorithms are part of the python Sklearn library: in order to implement and train the algorithm we have to first parse the data inside the `train_dataset.json` file, then we need to sort all the data by category in a CSV file, in order to train and test our model.

4.1 Parsing the Train Set

As we said, the jsonl needs to be parsed and converted in a CSV in order to train our model, this can be achieved by parsing the file line-by-line:

```
if not os.path.isfile(input_path):
    print("File path {} does not exist. Exiting...".format(filepath))

inputFile = open(input_path, 'r');
outputFile = open(output_path, 'w');

line=inputFile.readline()
output = csv.writer(outputFile,delimiter='\\t')

while line:
    data = json.loads(line)
    instr=""
    count=0
    for command in data["instructions"]:
        if command.split(" ")[0][0]=='m':
            count+=1
            instr=instr+command+" "
    output.writerow([ instr,count,data["opt"],data["compiler"], len(instr.split(" "))])
    line = inputFile.readline()

inputFile.close() #close the input file
```

```
outputFile.close()#close the output csv file
```

Then we can create our db of instructions to feed to our model:

```
filename = 'train_datasetCSV.csv'
db = pd.read_csv(filename, sep='\t', header=None,
names=['instructions','mov', 'opt','compiler','LOC'],
dtype={'mov': np.int32, "LOC": np.int32})
```

4.2 Vectorizing the Data

In order to efficiently train our models is not enough to feed them the raw files, we need to vectorize them. We can apply various strategies to achieve it: for example we can use the Count Vectorizer method that assign to each word in the vocabulary a token to feed to the machine learning algorithm.

Another vectorizer known as Tf-idf Vectorizer, count the occurrences of each words inside the documents and lastly we have Count Vectorizer combined with n-grams, that check chunks of 1 to n words, instead of building the vocabulary word-by-word.

Here is the implementation:

```
def countVectorizer(db):

    start_time = time.time()

    x=[]
    y=[]
    x_new=[]

    for ins in db.instructions:
        temp=ins.split(",")
        s=""
        for cmd in temp:
            s=s+cmd.split(" ")[0]+" "
        x_new.append(s)

    vectorizer = CountVectorizer(stop_words=None)
    x = vectorizer.fit_transform(x_new)
    y = db.opt

    X_train, X_test, y_train, y_test = train_test_split(x, y,
                                                         test_size=0.3)

    print("Count vectorizer:")
    print("Train: %d - Test: %d" %(X_train.shape[0],X_test.shape[0]))
    elapsed_time = time.time() - start_time
    print("Vectorizer elapsed Time: %.3fs" %elapsed_time)
```

```

        print("\n")
        return X_train, X_test, y_train, y_test

countVectorizer(db)

def tfidfVectorizer(db):

    start_time = time.time()

    x=[]
    y=[]
    x_new=[]

    for ins in db.instructions:
        temp=ins.split(",")
        s=""
        for cmd in temp:
            s=s+cmd.split(" ")[0]+" "
        x_new.append(s)

    vectorizer=TfidfVectorizer()
    x = vectorizer.fit_transform(x_new)
    y = db.opt

    X_train, X_test, y_train, y_test = train_test_split(x, y,
                                                         test_size=0.3)

    print("tfidfVectorizer:")
    print("TrainSet: %d - TestSet: %d" %(X_train.shape[0],X_test.shape[0]))
    elapsed_time = time.time() - start_time
    print("TfidfVectorizer elapsed Time: %.3fs" %elapsed_time)
    print("\n")

    return X_train, X_test, y_train, y_test

tfidfVectorizer(db)

def vectorizer2Gram(db):

    start_time = time.time()

    vectorizer = CountVectorizer(ngram_range=(2,2)) # multinomial
    x = vectorizer.fit_transform(db.instructions)
    y = db.opt

    X_train, X_test, y_train, y_test = train_test_split(x, y,
                                                         test_size=0.3)

```

```

print("vectorizer2Gram:")
print("TrainSet: %d - TestSet: %d" %(X_train.shape[0],X_test.shape[0]))
elapsed_time = time.time() - start_time
print("vectorize2Gram elapsed Time: %.3fs" %elapsed_time)
print("\n")

return X_train, X_test, y_train, y_test

vectorizer2Gram(db)

```

4.3 Splitting the DataSet

Taking a closer look to the execution of a vectorization, we can see that the DataSet is not evenly splitted, that's because giving more data to the train set makes the learning algorithm more precise:

```

tfidfVectorizer:
TrainSet: 21000 - TestSet: 9000
TfidfVectorizer elapsed Time: 11.012s

```

4.4 Creating the Model and Processing the Data

Now that we have the vectorized data ready, we can feed them to a learning algorithm, here is the implementation of the most efficient combination between vectorization and learning algorithm:

```

def DecisionTreeOPT():

    start_time = time.time()

    print("DecisionTree - countVectorizer")

    xTrain, xTest, yTrain, yTest=countVectorizer(db)

    model = DecisionTreeClassifier(random_state=0).fit(xTrain, yTrain)
    yPred = model.predict(xTest)

    print("confusion Matrix of DT/ cVectorizer:\n")
    print(confusion_matrix(yTest, yPred))
    print("\n")
    print("classificationReport Matrix of DT/ cVectorizer:\n")
    print(classification_report(yTest, yPred))

    plotPrecisionRecallOPT(yTest,yPred,"DecisionTree - countVectorizer")

```

```

elapsed_time = time.time() - start_time
print("elapsed Time: %.3fs" %elapsed_time)
print("-----")

print("DecisionTree - tfidfVectorizer")

xTrain, xTest, yTrain, yTest=tfidfVectorizer(db)

model = DecisionTreeClassifier(random_state=0).fit(xTrain, yTrain)
yPred = model.predict(xTest)

print("confusion Matrix of DT/ tfidfVectorizer:\n")
print(confusion_matrix(yTest, yPred))
print("\n")
print("classificationReport Matrix of DT/ tfidfVectorizer:\n")
print(classification_report(yTest, yPred))

plotPrecisionRecallOPT(yTest,yPred,"DecisionTree - tfidfVectorizer")

elapsed_time = time.time() - start_time
print("elapsed Time: %.3fs" %elapsed_time)
print("-----")

print("DecisionTree - vectorizer2Gram")

xTrain, xTest, yTrain, yTest=vectorizer2Gram(db)

model = DecisionTreeClassifier(random_state=0).fit(xTrain, yTrain)
yPred = model.predict(xTest)

print("confusion Matrix of DT/ vectorizer2Gram:\n")
print(confusion_matrix(yTest, yPred))
print("\n")
print("classificationReport Matrix of DT/ vectorizer2Gram:\n")
print(classification_report(yTest, yPred))

plotPrecisionRecallOPT(yTest,yPred,"DecisionTree - vectorizer2Gram")

elapsed_time = time.time() - start_time
print("elapsed Time: %.3fs" %elapsed_time)
print("-----")

DecisionTreeOPT()

```


5 Evaluating the Performances

The most critical aspect of every learning algorithm while approaching a text classification problem, is of course the ability to correctly predict the category where a text or a word will fit in: this paragraph will describe all the criteria used to evaluate all the different combination of vectorization and learning algorithms:

5.1 Precision

The precision of the learning algorithm measures the proportion of the positives that were correctly identified, as the formula suggests:

$$Precision = \frac{TP}{TP+FP}.$$

The precision represents the ratio between true positive (positive classified correctly by the algorithm) and the sum of true positive and false positive (positive wrongly classified as such).

5.2 Recall

The recall measures the percentage of actual positive correctly predicted by the algorithm, so we can define recall as follows:

$$Recall = \frac{TP}{TP+FN}.$$

To measure how good a learning algorithm is, we have to measure both the recall and the accuracy, usually incrementing one comes at the expense of the other.

5.3 F1 Score

The F1 score is the indicator of a test accuracy: it considers both precision and recall to compute the score:

$$F1 = 2 \times \frac{precision \times recall}{precision + recall}.$$

So here is a classification report (obtained by the command `classification_report(yTest, yPred)`) of the best combination(in terms of performances) of the vectorizer2gram ad Decision Tree learning algorithm (Binary classifier scenario):

```
classificationReport Matrix of DT/ vectorizer2Gram:
```

	precision	recall	f1-score	support
H	0.80	0.80	0.80	3609
L	0.86	0.86	0.86	5391
accuracy			0.84	9000
macro avg	0.83	0.83	0.83	9000
weighted avg	0.84	0.84	0.84	9000

5.4 The Precision-Recall curve

Beside looking to the precision and recall values themselves, we can plot an useful graph that helps evaluate performances between learning algorithms: it's the precision-recall curve.

The PR curve plot the precision on the X axis and the recall on the Y axis, it measures the trade-off between the true positive rate and the positive predictive value for a predictive model, using different probability thresholds, here are two executions for the binary classification problem(High and Low instructions) for a good and a bad learning algorithm:

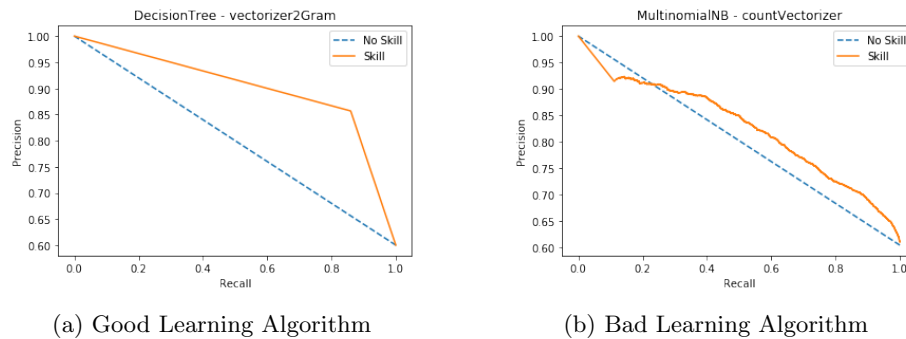


Figure 1: PR curve of two learning algorithm

The difference lies in the curve being below or over the no-skill line, basically the algorithm in figure(b) has less probability of correctly categorizing a word than randomly putting in a category.

Here is the equivalent curve for a multiclass learning algorithm:

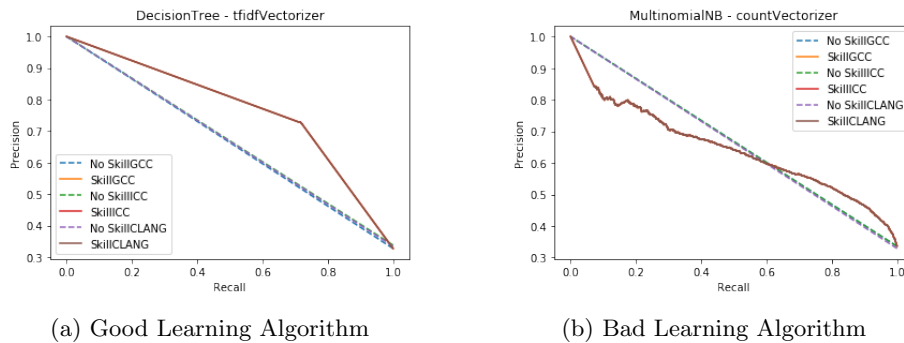


Figure 2: PR curve of two multiclass learning algorithm

6 BlindTest

Once we evaluated the best combination of vectorization and learning algorithm for both the binary classification problem and multiclass, we need to evaluate how they would react to a different set of data once they are trained.

6.1 Binary Classification Problem

So the first thing to do is to parse the `test_dataset_blind.jsonl` :

```
if not os.path.isfile(input_path_blind):
    print("File path {} does not exist. Exiting...".format(filepath))

inputFile = open(input_path_blind, 'r');
line=inputFile.readline()

blind_test_data = [];
while line:
    data = json.loads(line)
    instr=""
    for command in data["instructions"]:
        instr = instr + command+", "

    blind_test_data.append(instr)
    line = inputFile.readline()

print(len(blind_test_data))
```

Now it's time to train back the most efficient algorithms (binary and multiclass):

```
model,vectorize = DecisionTreeOPT() #training with csv file
```

```
x_blind_test = vectorize.transform(blind_test_data)
y_pred = model.predict(x_blind_test)
```

```
low_instr=0.0
high_instr=0.0
```

```
for i in y_pred:
    if(i=="L"):
        low_instr+=1;
    else:
        high_instr+=1;
```

```
db_length = len(y_pred);
```

```
lPerc = (low_instr/db_length)*100;
hPerc = (high_instr/db_length)*100;
```

```
print("LOW=%.6f HIGH=%.6f" %(lPerc,hPerc));
```

classificationReport Matrix of DT/ vectorizer2Gram:

	precision	recall	f1-score	support
H	0.79	0.77	0.78	3636
L	0.85	0.86	0.86	5364
accuracy			0.83	9000
macro avg	0.82	0.82	0.82	9000
weighted avg	0.83	0.83	0.83	9000

```
elapsed Time: 155.786s
```

So here is the output for the reading of just the CSV file:

```
LOW=59.746667 HIGH=40.253333
```

And here is the percentage categorizing the BlindSet:

```
LOW=58.866667 HIGH=41.133333
```

So the algorithm just misclassified only a small fraction of LOW and HIGH instructions.

6.2 Multiclass Classification Problem

Again, we have to do the same thing for the multiclass learning problem:

```
model,vectorize = DecisionTreeCompiler()

x_blind_test = vectorize.transform(blind_test_data)
y_pred = model.predict(x_blind_test)

gcc_instr=0.0
icc_instr=0.0
clang_instr =0.0

for i in y_pred:
    if(i=="gcc"):
        gcc_instr+=1;
    elif(i=="icc"):
        icc_instr+=1;
    else:
        clang_instr+=1;

db_length = len(y_pred);

gccPerc = (gcc_instr/db_length)*100;
iccPerc = (icc_instr/db_length)*100;
clangPerc = (clang_instr/db_length)*100;

print("GCC=%.6f ICC=%.6f CLANG=%.6f" %(gccPerc ,iccPerc,clangPerc));

classificationReport Matrix of DT/ vectorizer2Gram:

              precision    recall  f1-score   support

 clang         0.92         0.91         0.91         3001
   gcc         0.90         0.91         0.90         2978
   icc         0.93         0.92         0.92         3021

 accuracy                   0.91         9000
 macro avg         0.91         0.91         0.91         9000
weighted avg         0.91         0.91         0.91         9000

elapsed Time: 116.865s
```

The original content of the DataSet were three equally divided set of GCC ,
ICC and CLANG set of instructions:

GCC=33.333333 ICC=33.333333 CLANG=33.333333

And here is the output of the learning algorithm using the blind data set:

GCC=33.866667 ICC=33.366667 CLANG=32.766667

Again the classifier just misclassified only a small fraction of the instructions, those result are compatible both with the the 0.91 accuracy of the report matrix and the precision-recall curve.

7 References

The Scikit-learn library for Python: <https://scikit-learn.org/stable/>

The Precision-Recall curve in Python using Scikit-learn: https://scikit-learn.org/stable/auto_examples/model_selection/plot_precision_recall.html